

Report of Lab 3

Yuxuan Duan 516030910573

1 Design Decisions

1.1 About the Operators

Considering that an operator may not actually do the operation, or may be closed & reopened or rewound, I chose to do necessary operations only at the first time *fetchNext()*. It may seem unnatural but it saves time.

1.2 About HashEquiJoin

1. When the first time *fetchNext()* is called, a `HashMap(child2Field.hashCode() -> ArrayList<Tuple>)` is created. And `child2` is sequentially scanned and appended to the corresponding `ArrayList` in the `HashMap`. The code above will only be run once.
2. Every time *fetchNext()* is called, `child1Field.hashCode()` is calculated and is used to index the `ArrayList` in `HashMap`. The iterator takes the next tuple in the `ArrayList`, uses `Predicate` to judge the two fields are indeed equal, and finally makes a new tuple and returns it.
3. Each time an `ArrayList` is indexed and the first eligible tuple is returned, the rest of the tuples in that `ArrayList` are all possibly eligible. Also, all the possibly eligible tuples are in that `ArrayList`. So until we have finished scanning that `ArrayList`, we can simply store that `ArrayList`, mark *fetching* = *true*, and continue scanning (directly starting Step 3) that `ArrayList` in the next few times we call *fetchNext()*. After we finish scanning, mark *fetching* = *false*, and we will head back to Step 2 next time we call *fetchNext()*. `Predicate` is also needed in Step 3.

Note: Using `Predicate` in Step 2 and 3 is necessary. Different fields may have a same hash code, so we use `Predicate` to further ensure the two fields are the same, even after we have known that their hash codes are the same.

1.3 About Join Operator Choice

As we currently haven't implemented the optimizer, so I added a special judge in the `Join` class. When a `Join` class is created with *op* == *Predicate.Op.EQUAL*, operations of `HashEquiJoin` will be done if *fetchNext()* is called.

In the future labs when I implement the optimizer, I will remove the above usage.

2 API Changes

2.1 `BufferPool.replacePage()`

Sometimes we need to create a new page without getting it through the `BufferPool`, edit it and write it back (e.g. Tests in `BufferPoolWriteTest`). If we try to mark a page dirty which is not in the `BufferPool`, `flushPage()` will actually do nothing. So we need a method to get a page into the `BufferPool` directly from a existing page in the memory.

`BufferPool.replacePage()`:

- input a existing page in the memory
- if this page (of an old version) is already in the `BufferPool`, replace the old page in the `BufferPool` with the new page.
- if this page is not in the `BufferPool`, simply add it.

2.2 `TupleDesc.setFieldName()`

It provides a convenient method to edit the name of a field and is used when an `Aggregate` class edits the `groupByFieldName` of the `TupleDesc` of its `returnIterator` according to the `groupByFieldName` of its `childIterator`. (`IntegerAggregator` and `StringAggregator` cannot do this because they do not know the `groupByFieldName` until `mergeTupleIntoGroup()` is called.)

3 Missing or Incomplete Elements

There was no more missing or incomplete element than those in Lab 2 to the best of my knowledge.

4 Time Spent and Difficulties

I roughly spent 20 hours for Lab 3. In this lab, I spent much time debugging and some of the bugs are in the classes I implemented in Lab 1 or 2. I think this happened because now I'm confronted with far more complicated tasks, and some special situations were not thought over in the previous labs. I still have some minor difficulties in the structure of `SimpleDB` yet they are not as confusing as they were in Lab 1 and 2.